

CS 372 Lecture #11

The Application Layer:

- More HTTP
- Cookies
- Caching

Note: Many of the lecture slides are based on presentations that accompany *Computer Networking: A Top Down Approach*, 6th edition, by Jim Kurose & Keith Ross, Addison-Wesley, 2013.

Client-server state: cookies

Many major Web sites use cookies

Four components:

- 1) cookie header line of HTTP *response* message
- 2) cookie header line in HTTP *request* message
- 3) cookie file kept on user's host, managed by user's browser
- 4) back-end database at Web site

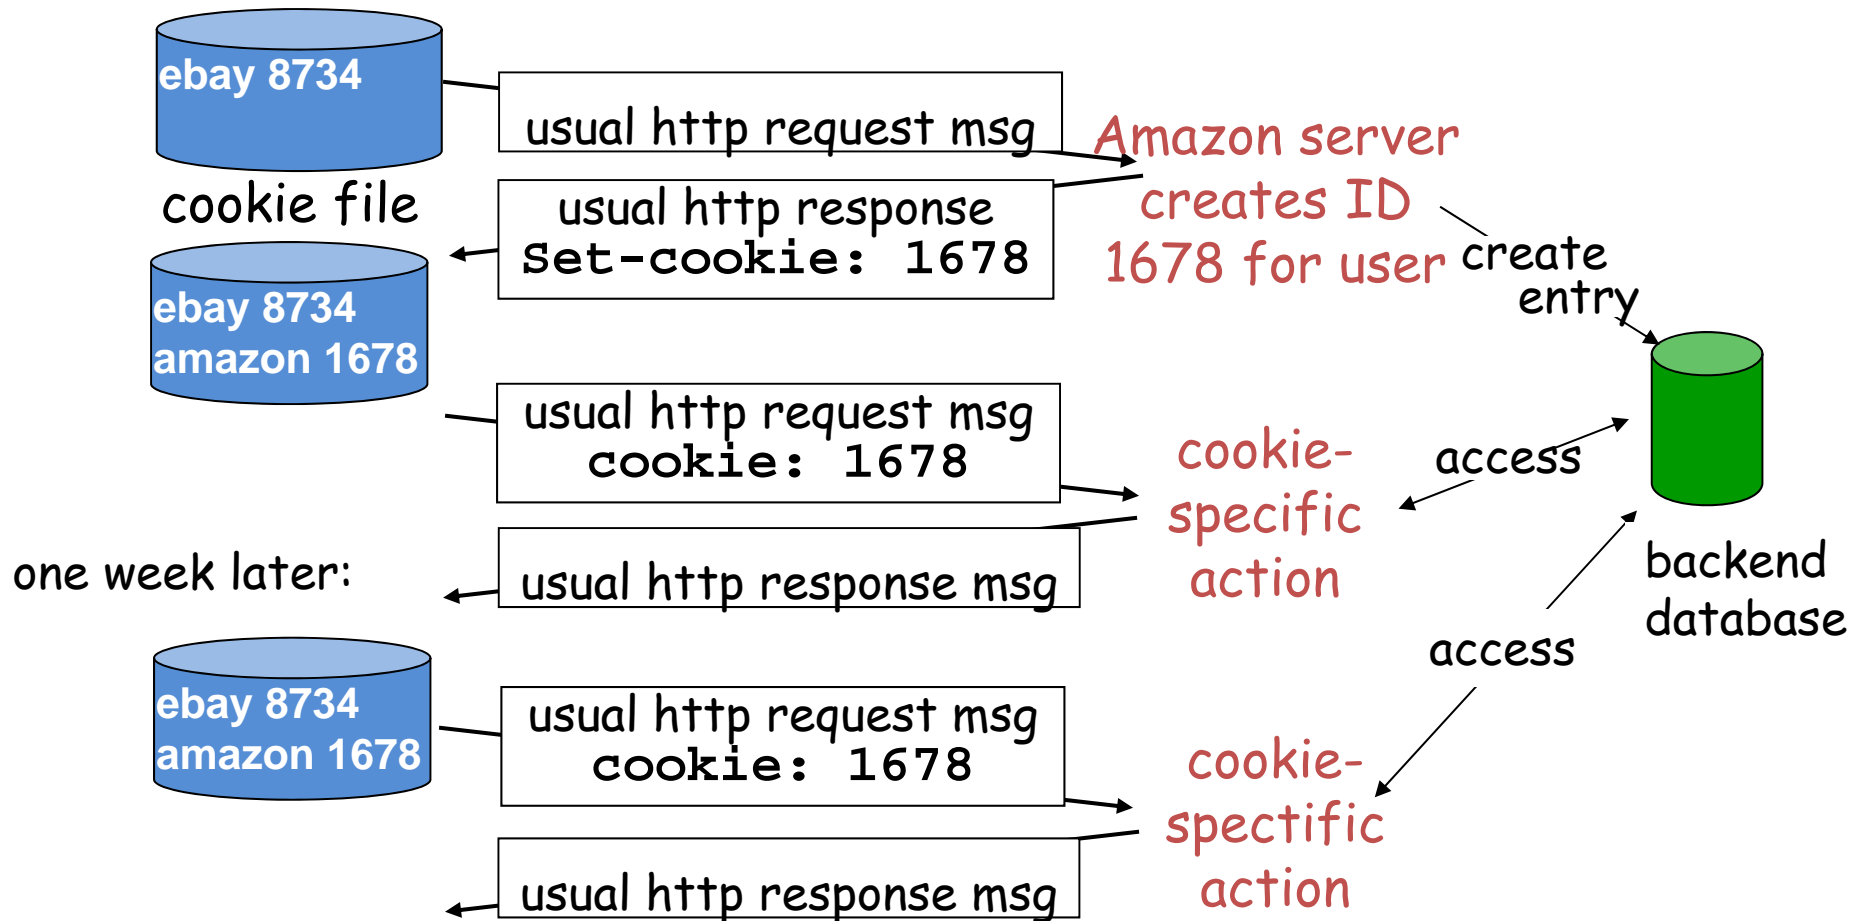
Example:

- User visits specific e-commerce site for first time
- when initial HTTP requests arrives at site, site creates:
 - unique ID
 - entry in backend database for ID

Cookies: keeping “state”

client

server



Cookies (cont.)

What cookies can provide:

- authorization
- shopping carts
- recommendations
- user session state (Web e-mail)


Cookies and privacy: aside

- cookies permit sites to learn a lot about you
- you may be giving your name and e-mail to sites

Web caches (proxy server)

Goal: satisfy client request without involving origin server

- User's browser sends all HTTP requests to cache
 - if object in cache: cache returns object
 - else cache requests object from origin server, then returns object to client



More about Web caching

- Cache acts as both client and server
- Typically cache is installed by ISP (university, company, residential ISP)
- Cached objects have “expiration” date/time

Why Web caching?

- reduce response time for client request
- reduce traffic on an institution’s access link.
- enables “poor” providers to effectively deliver content


Example (no caching)

Assumptions

- average object size = 100K bits
- average request rate from institution's browsers to origin servers = 15 requests per second
- delay from institutional router to any origin server and back to router = 2 seconds

Consequences

- utilization on LAN = 15%
- utilization on access link = ~100%



total average delay = Internet delay + access delay + LAN delay
= 2 seconds + minutes + milliseconds


Example (caching)

Same assumptions

- ... but with caching
- Suppose cache hit rate is 0.4
 - i.e. only 60% of requests go to origin servers

Consequences

- 40% of requests will be satisfied almost immediately
- utilization of access link is reduced to 60%, resulting in negligible delays (say 10 ms) because of no congestion



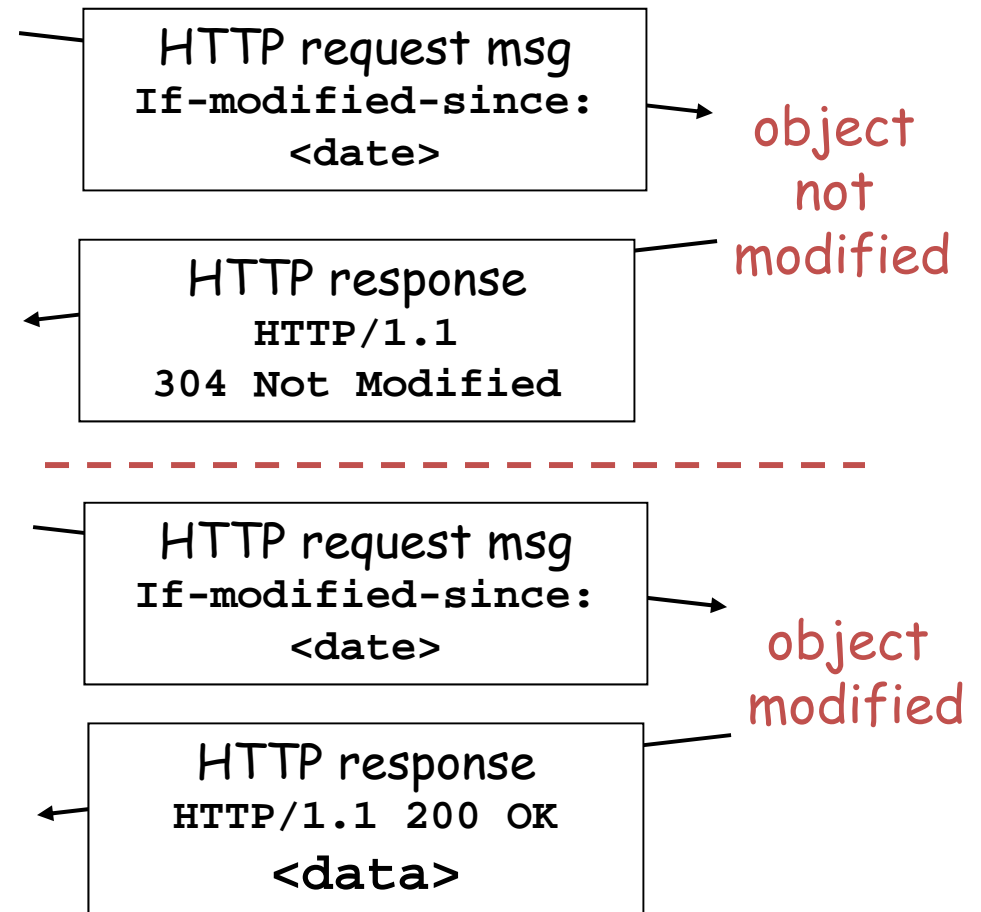
$$\begin{aligned} \text{total average delay} &= \text{Internet delay} + \text{access delay} + \text{LAN delay} \\ &= 0.6 \cdot (2.0) \text{ seconds} + 0.4 \cdot 10 \text{ milliseconds} < 1.4 \text{ seconds} \end{aligned}$$

Conditional GET

- **Goal:** don't send object if cache has up-to-date cached version
- cache: specify date of cached copy in HTTP request
If-modified-since:
<date>
- server: response contains no object if cached copy is up-to-date:
HTTP/1.1 304 Not Modified

cache

server



- Definitions
 - Cookie
 - Caching
- HTTP
 - Conditional GET